

AIG130 – Project 1

AI-Powered Hyper-Personalized Shopping Assistant

A Machine Learning Solution for Retail / E-Commerce

Cloud Platform: Google Cloud Platform (AutoML & Vertex AI)

Sector: Retail / E-Commerce

Course: AIG130 | Group Submission

Group-04

Srimedha Jaggavarapu

Alexandre Papouchine

AnuSri Saini

Mohammad Salarinejad

Chao Chen

1. Introduction

The modern e-commerce landscape is saturated with millions of products, making discovery increasingly difficult for consumers. While basic recommendation systems rely primarily on collaborative filtering based on past purchases and are largely reactive rather than proactive.

This project proposes an AI-Powered Hyper-Personalized Shopping Assistant which is a real-time, conversational ML system that combines natural language understanding, computer vision, behavioral analytics, and contextual awareness to deliver truly individualized shopping experiences. Unlike existing solutions, this assistant:

- Understands natural language queries.
- Adapts recommendations in real time based on current browsing session context
- Incorporates multi-modal inputs: text, images uploaded by the user, and voice
- Learns continuously from micro-interactions (hover time, scroll depth, wish list additions)
- Factors in external context: weather, local events, trending styles, inventory levels

The solution will be designed, implemented, and deployed on Google Cloud Platform (GCP), leveraging Vertex AI, AutoML and related services. This report covers the problem definition, platform analysis, full ML pipeline, implementation details, and evaluation of the solution's impact.

2. Problem Identification

2.1 Problem Statement

Current e-commerce recommendation engines suffer from three critical limitations:

- Cold Start Problem: New users receive irrelevant or generic recommendations due to lack of purchase history.
- Static Context: Recommendations ignore real-world context such as weather, upcoming occasions, or current trends.
- Passive Interaction: Users must search and filter manually; the system does not proactively engage or guide them.

These limitations result in high cart abandonment rates (estimated at 70%+ industry-wide), poor product discoverability, and lost revenue. The opportunity is significant: McKinsey estimates that 35% of Amazon's revenue is generated by its recommendation engine, yet most mid-to-large retailers lack comparable technology.

2.2 Proposed Solution

The AI-Powered Hyper-Personalized Shopping Assistant addresses these gaps through a multi-model ML architecture:

Component	Technology	Purpose
Conversational AI	Vertex AI + Gemini API	Natural language query understanding and dialogue management
Recommendation Engine	AutoML Tables + Collaborative Filtering	Product ranking and personalization
Visual Search	AutoML Vision	Match products from user-uploaded images
Sentiment Analysis	Natural Language API	Analyze product reviews and user feedback
Context Engine	Custom model on Vertex AI	Integrate weather, events, trends, and inventory
Behavioral Analytics	BigQuery ML	Real-time micro-interaction processing

2.3 Novelty & Innovation

This solution does not exist as a unified, commercially deployed system for mid-market retailers. Key innovations include:

- Multi-modal fusion: combining text, image, and behavioral signals in a single inference pipeline
- Real-time context injection: live weather API + event calendar + inventory feeds update recommendations dynamically
- Conversational shopping: a chat/voice interface that guides users through complex purchase decisions
- Continuous online learning: the model updates from every session without full retraining

3. Cloud Platform Analysis & Justification

3.1 Platform Comparison

We evaluated three major cloud platforms — AWS, Google Cloud Platform (GCP), and Microsoft Azure — across six criteria relevant to our ML solution.

Criteria	AWS SageMaker	Google Cloud AutoML / Vertex AI	Azure ML Studio
Scalability	Excellent – auto-scaling with SageMaker Endpoints	Excellent – Vertex AI scales globally	Very Good – AKS-based scaling
Cost	Moderate – pay-per-instance; can be expensive at scale	Competitive – serverless options reduce idle cost	Moderate – hybrid/reserved options available
Accessibility	Broad – wide documentation and community	Very Good – strong Python SDK, Colab integration	Good – tight Microsoft ecosystem integration

Criteria	AWS SageMaker	Google Cloud AutoML / Vertex AI	Azure ML Studio
Computational Power	High – P3/P4 GPU instances	High – TPUs available for large models	High – NC-series GPU instances
Ease of Use	Moderate – powerful but complex configuration	High – AutoML requires minimal ML expertise	High – drag-and-drop Designer UI
ML Services Breadth	Very High – comprehensive ecosystem	High – strong NLP, Vision, Tables AutoML	High – strong Azure OpenAI integration

3.2 Justification for Google Cloud Platform

GCP / Vertex AI was selected for the following reasons:


- AutoML Tables & Vision: Directly applicable to our recommendation engine and visual search components with minimal custom code required.
- Vertex AI Unified Platform: Allows us to manage data, train models, deploy endpoints, and monitor predictions from a single interface, reducing operational complexity.
- BigQuery ML: Enables ML directly on structured behavioral data without data movement is critical for real-time session analytics.
- Gemini API Access: GCP provides native access to Gemini models for the conversational AI component.
- Free Tier & Student Credits: GCP offers \$300 in free credits, making it the most accessible option for a student group.
- TPU Availability: For future scaling of deep learning models, GCP's TPUs provide a significant performance advantage.

3.3 Limitations of GCP

- Vendor Lock-in: Heavy use of GCP-native services (Vertex AI, BigQuery, AutoML) makes migration to other platforms costly.
- Learning Curve: Vertex AI pipelines require familiarity with Kubeflow and GCP IAM, which can be complex for beginners.
- Cost at Scale: AutoML prediction costs can become significant at high query volumes without careful budget management.

4. Comprehensive ML Pipeline Implementation

The ML pipeline follows a full MLOps lifecycle, covering data collection through continuous monitoring. Each stage is designed around GCP services. Our group implemented all four stages of the pipeline.



Stage 1: Data Collection & Preprocessing

We used the publicly available Amazon Product Reviews dataset (Electronics category, ~1.7M reviews) as a proxy for real user-behavioral data. The preprocessing pipeline was implemented as follows:

- Data ingestion: Raw CSV files uploaded to Cloud Storage; loaded into BigQuery using bq load
- Feature engineering: BigQuery SQL + Python (Pandas) used to generate user-item interaction matrices and session features
- AutoML Tables dataset: Preprocessed data registered in Vertex AI Datasets for training
- Tool used: Google Colab (connected to GCP) for exploratory data analysis and feature engineering scripts

Dataset is available in BigQuery:

<https://console.cloud.google.com/bigquery?project=project-1-489519&ws=!1m0>

Stage 2: Model Training

We trained the Product Ranking sub-model using AutoML Tables on the preprocessed interaction dataset:

- Target variable: binary 'purchased' (1) vs 'browsed only' (0)
- Features: user_id embedding, product_category, price_tier, session_duration, time_of_day, day_of_week
- Training budget: 1 compute hour (within free tier limits)
- Best model: AutoML selected a Gradient Boosted Trees ensemble with AUC-ROC = 0.87 on validation set

We can monitor the training progress in :

<https://console.cloud.google.com/vertex-ai/training/training-pipelines?project=project-1-489519>

We can check the metrics of the trained model at:

<https://console.cloud.google.com/vertex-ai/models/locations/us-west1/models/7732165988753342464/versions/1/evaluations?project=project-1-489519>

Stage 3: Deployment (API/Web App)

The trained AutoML Tables model was deployed to a Vertex AI Endpoint and wrapped in a Cloud Run API:

- Vertex AI Endpoint: deployed with 1 node (n1-standard-4), auto-scaling enabled
- Cloud Run API: Python (FastAPI) service that receives user context, calls Vertex AI prediction, and returns ranked product list
- Frontend demo: Simple React.js chat interface demonstrating conversational product search
- The API returns top 5 recommended product IDs with confidence scores in < 150ms average latency

We can see the trained & deployed model status in the Model Registry at :

<https://console.cloud.google.com/vertex-ai/models?project=project-1-489519>

We can verify whether the shopping assistant is live in Cloud Run using the link:

<https://shopping-assistant-api-zueyk36pqa-uw.a.run.app/docs#/>

Stage 4: Monitoring & MLOps

- Vertex AI Model Monitoring enabled with 1-hour monitoring window
- Feature drift alerts configured for key features (price_tier, session_duration)
- Cloud Logging dashboard created to visualize API request volume and error rates
- Automated retraining trigger: Cloud Scheduler + Cloud Functions checks drift metrics daily

We can also check the monitoring status at:

<https://console.cloud.google.com/monitoring/settings/metric-scope?project=project-1-489519>

4.2 Challenges & Solutions

Challenge	Solution
AutoML training quota limits on free tier	Reduced training budget to 1 hour; used smaller dataset sample (100K rows)
Cold start for new users	Implemented popularity-based fallback; new users receive trending items by category
API latency spikes under load	Added Redis caching layer; pre-computed recommendations for top 1000 users
Data privacy (user behavioral data)	Anonymized all user IDs; complied with GCP's data processing addendum
Model monitoring setup complexity	Used Vertex AI's built-in monitoring templates; set conservative drift thresholds

5. Solution Evaluation

5.1 Model Performance

Metric	Baseline (Popularity)	Our AutoML Model	Improvement
AUC-ROC	0.71	0.87	+22.5%
Precision@5	0.18	0.31	+72.2%
NDCG@10	0.22	0.39	+77.3%
Mean Reciprocal Rank	0.14	0.28	+100.0%
Avg. Inference Latency	N/A	148ms	< 200ms target ✓

5.2 Economic Impact

- Cart Abandonment Reduction: Hyper-personalized recommendations are estimated to reduce abandonment by 15-20%, directly increasing conversion rates.
- Revenue Uplift: Industry studies suggest personalization engines drive 10-30% incremental revenue. For a mid-size retailer with \$50M annual revenue, this represents \$5M-\$15M additional revenue.
- Customer Lifetime Value: Improved product-fit matching increases repeat purchase rates and customer satisfaction scores (NPS).
- Operational Cost Savings: AutoML eliminates the need for a dedicated ML engineering team for model development, reducing costs by an estimated \$200K-\$500K/year.

5.3 Broader Implications

- Privacy: The system collects granular behavioral data. GDPR/CCPA compliance is essential; users must opt-in to personalization and have the right to delete their data profile.
- Algorithmic Bias: Recommendation systems can create 'filter bubbles' and disadvantage lower-margin products. Regular fairness audits are required.
- Environmental Impact: Large-scale ML inference has a carbon footprint. GCP's carbon-neutral infrastructure partially mitigates this; model compression techniques (distillation, quantization) should be applied.
- Small Business Impact: This technology, when accessible via GCP's pay-per-use model, can democratize personalization for SME retailers who previously could not afford such systems.

6. Conclusion

This project successfully designed, implemented, and evaluated an AI-Powered Hyper-Personalized Shopping Assistant for the retail/e-commerce sector is a solution that addresses critical gaps in today's recommendation landscape by combining conversational AI, multi-modal inputs, real-time context, and continuous learning.

We selected Google Cloud Platform as our implementation environment, leveraging Vertex AI, AutoML Tables, Cloud Run, and BigQuery ML. The full ML pipeline - from data collection through monitoring was implemented and validated, with our AutoML model achieving a 77% improvement in NDCG@10 over the popularity-based baseline.

The solution demonstrates that modern cloud ML platforms have democratized access to state-of-the-art recommendation technology. The combination of GCP's managed services, AutoML's low-code training, and Vertex AI's unified MLOps platform allowed a student team to build a production-grade system within academic constraints.

Future work should focus on:

- (1) integrating the Gemini API for full conversational capability
- (2) implementing the visual search pipeline with AutoML Vision
- (3) deploying a real A/B test with live traffic
- (4) conducting a comprehensive fairness and bias audit across demographic segments.

The economic analysis indicates significant ROI potential, with personalization engines driving 10-30% revenue uplift for retailers making this not only a technically compelling solution but a commercially valuable one.

References

1. Google Cloud. (2024). Vertex AI Documentation. <https://cloud.google.com/vertex-ai/docs>
2. Google Cloud. (2024). AutoML Tables Documentation. <https://cloud.google.com/automl-tables/docs>
3. McKinsey & Company. (2013). How retailers can keep up with consumers. McKinsey Quarterly.
4. Baymard Institute. (2024). Cart Abandonment Rate Statistics. <https://baymard.com/lists/cart-abandonment-rate>
5. Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep Learning. MIT Press.
6. Amazon Web Services. (2024). Amazon SageMaker Documentation. <https://docs.aws.amazon.com/sagemaker/>
7. Microsoft Azure. (2024). Azure Machine Learning Documentation. <https://docs.microsoft.com/en-us/azure/machine-learning/>
8. He, X., et al. (2017). Neural Collaborative Filtering. WWW 2017.
9. Sculley, D., et al. (2015). Hidden Technical Debt in Machine Learning Systems. NIPS 2015.
10. European Commission. (2018). General Data Protection Regulation (GDPR). Official Journal of the EU.